

Taken from: [https://colab.research.google.com/drive/1E2RViy7xmor0mhqskZV14\\_NUj2jMpJz3?authuser=1#scrollTo=9AdpFVfy6ya9](https://colab.research.google.com/drive/1E2RViy7xmor0mhqskZV14_NUj2jMpJz3?authuser=1#scrollTo=9AdpFVfy6ya9)

```
#import libraries
import numpy as np
import time
from datetime import datetime
import pandas as pd
import random
import networkx as nx
import matplotlib.pyplot as plt

#define the shape of the environment (i.e., its states)
environment_rows = 11
environment_columns = 11

#Create a 3D numpy array to hold the current Q-values for each state and action pair: Q(s, a)
#The array contains 11 rows and 11 columns (to match the shape of the environment), as well as a third "action" dimension.
#The "action" dimension consists of 4 layers that will allow us to keep track of the Q-values for each possible action in
#each state (see next cell for a description of possible actions).
#The value of each (state, action) pair is initialized to 0.
q_m = np.zeros((environment_rows, environment_columns, 4))
# q_values = np.random.randint(5, size = (11, 11))
q_m
```

[illegible]

```
q_m[(5,3)]
```

```
➡ array([0., 0., 0., 0.])
```

## ▼ Actions

The actions that are available to the AI agent are to move the robot in one of four directions:

- Up
- Right
- Down
- Left

Obviously, the AI agent must learn to avoid driving into the item storage locations (e.g., shelves)!

```
#define actions
#numeric action codes: 0 = up, 1 = right, 2 = down, 3 = left
actions = ['up', 'right', 'down', 'left']
```

## ▼ Rewards

The last component of the environment that we need to define are the **rewards**.

To help the AI agent learn, each state (location) in the warehouse is assigned a reward value.

The agent may begin at any white square, but its goal is always the same: ***to maximize its total rewards!***

Negative rewards (i.e., **punishments**) are used for all states except the goal.

- This encourages the AI to identify the *shortest path* to the goal by *minimizing its punishments!*

To maximize its cumulative rewards (by minimizing its cumulative punishments), the AI agent will need find the shortest paths between the item packaging area (green square) and all of the other locations in the warehouse where the robot is allowed to travel (white squares). The agent will also need to learn to avoid crashing into any of the item storage locations (black squares)!

```
# #Create a 2D numpy array to hold the rewards for each state.
# #The array contains 11 rows and 11 columns (to match the shape of the environment), and each value is initialized to -100.
# rewards = np.full((environment_rows, environment_columns), -100.)
# rewards[0, 5] = 100. #set the reward for the packaging area (i.e., the goal) to 100

# #define aisle locations (i.e., white squares) for rows 1 through 9
# aisles = {} #store locations in a dictionary
# aisles[1] = [i for i in range(1, 10)]
# aisles[2] = [1, 7, 9]
# aisles[3] = [i for i in range(1, 8)]
# aisles[3].append(9)
# aisles[4] = [3, 7]
# aisles[5] = [i for i in range(11)]
# aisles[6] = [5]
# aisles[7] = [i for i in range(1, 10)]
# aisles[8] = [3, 7]
# aisles[9] = [i for i in range(11)]

# #set the rewards for all aisle locations (i.e., white squares)
# for row_index in range(1, 10):
#     for column_index in aisles[row_index]:
#         rewards[row_index, column_index] = -1.

# #print rewards matrix
# for row in rewards:
#     print(row)

# rewards
```

Alternate env

```

#Create a 2D numpy array to hold the rewards for each state.
#The array contains 11 rows and 11 columns (to match the shape of the environment), and each value is initialized to -100.
rewards = np.full((environment_rows, environment_columns), -100.)
# rewards[0, 5] = 100. #set the reward for the packaging area (i.e., the goal) to 100
rewards[1, 0] = 100. #set the reward for the packaging area (i.e., the goal) to 100
rewards[6, 0] = 40. # Set sub-optimal reward

#define aisle locations (i.e., white squares) for rows 1 through 9
aisles = {} #store locations in a dictionary
aisles[0] = [i for i in range(0, 9)]
aisles[1] = [8, 10]
aisles[2] = [i for i in range(0, 7)]
aisles[2].append(8)
aisles[2].append(10)
aisles[3] = [i for i in range(6, 11)]
aisles[4] = [i for i in range(1, 5)]
aisles[4].append(6)
aisles[5] = [2, 4, 6, 7, 8, 9]
aisles[6] = [2, 4, 6]
aisles[7] = [0, 1, 2, 4, 6, 7, 8, 9, 10]
aisles[8] = [4, 10]
aisles[9] = [i for i in range(11)]

#set the rewards for all aisle locations (i.e., white squares)
for row_index in range(0, 10):
    for column_index in aisles[row_index]:
        rewards[row_index, column_index] = -1.

# Add puddles (penalties)
rewards[2, 4] = -40
rewards[4, 3] = -20

# #print rewards matrix
# for row in rewards:
#     print(row)
pd.DataFrame(rewards)

```



	0	1	2	3	4	5	6	7	8	9	10
0	-1.0	-1.0	-1.0	-1.0	-1.0	-1.0	-1.0	-1.0	-1.0	-100.0	-100.0
1	100.0	-100.0	-100.0	-100.0	-100.0	-100.0	-100.0	-100.0	-1.0	-100.0	-1.0
2	-1.0	-1.0	-1.0	-1.0	-40.0	-1.0	-1.0	-100.0	-1.0	-100.0	-1.0
3	-100.0	-100.0	-100.0	-100.0	-100.0	-100.0	-1.0	-1.0	-1.0	-1.0	-1.0
4	-100.0	-1.0	-1.0	-20.0	-1.0	-100.0	-1.0	-100.0	-100.0	-100.0	-100.0
5	-100.0	-100.0	-1.0	-100.0	-1.0	-100.0	-1.0	-1.0	-1.0	-1.0	-100.0
6	40.0	-100.0	-1.0	-100.0	-1.0	-100.0	-1.0	-100.0	-100.0	-100.0	-100.0
7	-1.0	-1.0	-1.0	-100.0	-1.0	-100.0	-1.0	-1.0	-1.0	-1.0	-1.0
8	-100.0	-100.0	-100.0	-100.0	-1.0	-100.0	-100.0	-100.0	-100.0	-100.0	-1.0
9	-1.0	-1.0	-1.0	-1.0	-1.0	-1.0	-1.0	-1.0	-1.0	-1.0	-1.0
10	-100.0	-100.0	-100.0	-100.0	-100.0	-100.0	-100.0	-100.0	-100.0	-100.0	-100.0

```
rewards[8, 0]
```



```
-100.0
```

	0	1	2	3	4	5	6	7	8	9	10
0	-1	-1	-1	-1	-1	-1	-1	-1	-1	-100	-100
1	100	-100	-100	-100	-100	-100	-100	-100	-1	-100	-1
2	-1	-1	-1	-1	-40	-1	-1	-100	-1	-100	-1
3	-100	-100	-100	-100	-100	-100	-1	-1	-1	-1	-1
4	-100	-1	-1	-20	-1	-100	-1	-100	-100	-100	-100
5	-100	-100	-1	-100	-1	-100	-1	-1	-1	-1	-100
6	40	-100	-1	-100	-1	-100	-1	-100	-100	-100	-100
7	-1	-1	-1	-100	-1	-100	-1	-1	-1	-1	-1
8	-100	-100	-100	-100	-1	-100	-100	-100	-100	-100	-1
9	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
10	-100	-100	-100	-100	-100	-100	-100	-100	-100	-100	-100

## Train the Model

Our next task is for our AI agent to learn about its environment by implementing a Q-learning model. The learning process will follow these steps:

1. Choose a random, non-terminal state (white square) for the agent to begin this new episode.
2. Choose an action (move *up*, *right*, *down*, or *left*) for the current state. Actions will be chosen using an *epsilon greedy algorithm*. This algorithm will usually choose the most promising action for the AI agent, but it will occasionally choose a less promising option in order to encourage the agent to explore the environment.
3. Perform the chosen action, and transition to the next state (i.e., move to the next location).
4. Receive the reward for moving to the new state, and calculate the temporal difference.
5. Update the Q-value for the previous state and action pair.
6. If the new (current) state is a terminal state, go to #1. Else, go to #2.

This entire process will be repeated across 1000 episodes. This will provide the AI agent sufficient opportunity to learn the shortest paths between the item packaging area and all other locations in the warehouse where the robot is allowed to travel, while simultaneously avoiding crashing into any of the item storage locations!

```
# Number of white squares
sum(rewards == -1).sum()
```

59

## Define Helper Functions

```

#define a function that determines if the specified location is a terminal state
def is_terminal_state(current_row_index, current_column_index):
    """
    This function is to determine whether the current state is a state which will terminate the game or not.
    If this function returns True, then it will terminate the game.
    Landing on a white square continues the game, while landing on a black or the green square will terminate it.
    """
    if rewards[current_row_index, current_column_index] == -1.: # Falling on white
        return False
    else: # Falling on green or black
        return True

#define a function that will choose a random, non-terminal starting location
def get_starting_location(row = None, col = None):
    current_row_index = np.random.randint(environment_rows) if row is None else row
    current_column_index = np.random.randint(environment_columns) if col is None else col

    # # get a random row and column index
    # current_row_index = np.random.randint(environment_rows)
    # current_column_index = np.random.randint(environment_columns)

    # # Get specific row and column
    # current_row_index = 9
    # current_column_index = 10

    #continue choosing random row and column indexes until a non-terminal state is identified
    #(i.e., until the chosen state is a 'white square').
    while is_terminal_state(current_row_index, current_column_index):
        current_row_index = np.random.randint(environment_rows)
        current_column_index = np.random.randint(environment_columns)
    return current_row_index, current_column_index

# Epsilon greedy policy
def eps_greedy(current_row_index, current_column_index, epsilon):
    #if a randomly chosen value between 0 and 1 is less than epsilon,
    #then choose the most promising value from the Q-table for this state.
    if np.random.random() < epsilon:
        return np.random.randint(4)
    else: #choose a random action
        return np.argmax(q_m[current_row_index, current_column_index])

# Random policy
def random_pol(current_row_index, current_column_index):
    return np.random.randint(4)

# Boltzmann policy (ref: https://github.com/8Gitbrix/Reinforcement-Learning/blob/master/qlearn.py and https://automaticaddison.com/boltz-policy/)
def boltz_policy(current_row_index, current_column_index, tau):
    if tau > 0:
        p = np.array([q_m[current_row_index, current_column_index, x]/tau for x in range(4)], dtype=np.float128)
        prob_actions = np.exp(p) / np.sum(np.exp(p))
        cumulative_probability = 0.0
        choice = random.uniform(0,1)
        for a,pr in enumerate(prob_actions):
            cumulative_probability += pr
            if cumulative_probability > choice:
                return a
    else:
        return np.argmax(q_m[current_row_index, current_column_index])

# Transition function
def get_next_location(current_row_index, current_column_index, action_index):
    """
    This will instruct how the agent will move. As long as the agent is not near one of the four walls, it can move about freely.
    The moment it is against a wall, it will not be able to move beyond it, and the next location will be the same as the current location.
    """
    new_row_index = current_row_index
    new_column_index = current_column_index
    if actions[action_index] == 'up' and current_row_index > 0:
        new_row_index -= 1
    elif actions[action_index] == 'right' and current_column_index < environment_columns - 1:
        new_column_index += 1
    elif actions[action_index] == 'down' and current_row_index < environment_rows - 1:
        new_row_index += 1
    elif actions[action_index] == 'left' and current_column_index > 0:
        new_column_index -= 1
    return new_row_index, new_column_index

#Define a function that will get the shortest path between any location within the warehouse that
#the robot is allowed to travel and the item packaging location.
def get_shortest_path(start_row_index, start_column_index):

```

```

#return immediately if this is an invalid starting location
if is_terminal_state(start_row_index, start_column_index):
    return []
else: #if this is a 'legal' starting location
    current_row_index, current_column_index = start_row_index, start_column_index
    shortest_path = []
    shortest_path.append([current_row_index, current_column_index])
    #continue moving along the path until we reach the goal (i.e., the item packaging location)
    while not is_terminal_state(current_row_index, current_column_index):

        # Choose policy
        action_index = eps_greedy(current_row_index, current_column_index, .0) # Epsilon greedy policy
        # action_index = boltz_policy(row_index, column_index, tau) # Boltzmann policy
        # action_index = random_pol(row_index, column_index) # Random policy

        #move to the next location on the path, and add the new location to the list
        current_row_index, current_column_index = get_next_location(current_row_index, current_column_index, action_index)
        shortest_path.append([current_row_index, current_column_index])
        # print(current_row_index, current_column_index)
    return shortest_path

# #display a few shortest paths
# start = time.time()
# print(f"Optimal path finding started: {datetime.fromtimestamp(start)}")

# print(f"\nPath: {get_shortest_path(0, 8)} (Moves: {len(get_shortest_path(3, 9))})" #starting at row 3, column 9
# print(f"\nPath: {get_shortest_path(5, 9)} (Moves: {len(get_shortest_path(5, 0))})" #starting at row 5, column 0
# print(f"\nPath: {get_shortest_path(9, 10)} (Moves: {len(get_shortest_path(9, 5))})" #starting at row 9, column 5

# end = time.time()
# print(f"\nOptimal path finding ended: {datetime.fromtimestamp(end)}")
# elapsed_time = end - start
# print(f"Elapsed time: {np.round(elapsed_time/60, 3)} minutes\n")

```

We can see before training, in this particular starting positions, the agent takes very long to converge, hence we had to interrupt it half way.

After training however, the agent manages to get there instantly

## ✓ Train the AI Agent using Q-Learning

```

# #define training parameters
# epsilon = [0.1, 0.5, 0.9] # exploration
# max_epsilon = 1
# min_epsilon = 0.01
# epsilon_decay_rate = [0.1, 0.5, 0.9] # If you go beyond 1 it will make epsilon go to inf
# discount_factor = [0.1, 0.5, 0.9] #discount factor for future rewards
# learning_rate = [0.1, 0.5, 0.9] #the rate at which the AI agent should learn
# training_episodes = 1000
# tau = 5 # This is temperature for blottmann policy. Higher tau means more equal probability of taking an action. Lower tau is choosin

# param_search_q_values = []
# param_search_episodes = []
# param_search_steps = []
# param_search_params = []
# param_index = []

# for param_i, param in enumerate(epsilon):
#     q_values = []
#     episodes = []
#     steps = []
#     q_m = np.zeros((environment_rows, environment_columns, 4))

#     #run through training episodes
#     for episode in range(training_episodes):
#         # Define starting location. Keep it empty to randomise start location during training
#         row_index, column_index = get_starting_location()
#         step = 0

#         #continue taking actions (i.e., moving) until we reach a terminal state
#         # (i.e., until we reach the item packaging area or crash into an item storage location)
#         while not is_terminal_state(row_index, column_index):

#             # Choose the policy
#             action_index = eps_greedy(row_index, column_index, param) # Epsilon greedy policy
#             # action_index = boltz_policy(row_index, column_index, tau) # Boltzmann policy
#             # action_index = random_pol(row_index, column_index) # Random policy

#             # Decay the epsilon rate linearly
#             if param > 0.01:
#                 param *= epsilon_decay_rate[-1]

#             # # Temp decay
#             # if tau > 0.1:
#             #     tau -= 0.5

#             # # Decay the epsilon rate exponentially based on minimum and max values
#             # epsilon = min_epsilon/((max_epsilon-min_epsilon)*np.exp(-epsilon_decay_rate[-1])) # https://medium.com/@nancyjemi/level-up-un
#             # # epsilon = min_epsilon + (max_epsilon - min_epsilon) * np.exp(-1. * episode / epsilon_decay_rate[-1]) # Non-lienar decay

#             #perform the chosen action, and transition to the next state (i.e., move to the next location)
#             old_row_index, old_column_index = row_index, column_index #store the old row and column indexes
#             row_index, column_index = get_next_location(row_index, column_index, action_index)

#             # Bellman's equation:  $Q[s_{old},a] + \alpha*(R[s_{old},a] + \gamma*(\max(Q[s]) - Q[s_{old},a]))$ 
#             q_m[old_row_index, old_column_index, action_index] = q_m[old_row_index, old_column_index, action_index] + learning_rate[-1]*(re
#             print(f"\nAction taken: Go {actions[action_index]} (index = {action_index})\nMoving from {old_row_index, old_column_index} to {
#             print(f"Q value is: {q_m[old_row_index, old_column_index, action_index]}")
#             q_values.append(q_m[old_row_index, old_column_index, action_index])
#             episodes.append(episode)
#             step += 1
#             steps.append(step)

#             # print(f"\nQ matrix for going {actions[action_index]} is:\n {q_values[:, :, action_index]}")
#             # #receive the reward for moving to the new state, and calculate the temporal difference
#             # reward = rewards[row_index, column_index]
#             # old_q_value = q_values[old_row_index, old_column_index, action_index]
#             # temporal_difference = reward + (discount_factor * np.max(q_values[row_index, column_index])) - old_q_value

#             # #update the Q-value for the previous state and action pair
#             # new_q_value = old_q_value + (learning_rate * temporal_difference)
#             # q_values[old_row_index, old_column_index, action_index] = new_q_value
#             param_search_q_values.append(q_values)
#             param_search_episodes.append(episodes)
#             param_search_steps.append(steps)
#             param_search_params.append(param)
#             param_index.append(param_i)

#     print(f"\nTraining complete for {training_episodes} training episodes!")

```

numeric action codes: 0 = up, 1 = right, 2 = down, 3 = left

```
# pd.DataFrame(param_search_steps[100])
```

The reason that the q matrix is 3D, is that the 3rd dimension contains the q values for each action at any given point.

For example at (2,0), going up (action = 0) has a q value of 100, which makes sense because going up from (2,0) will win the game.

So then we look at the q matrix corresponding to action = 0, from the 4 actions. This means we access the q matrix that corresponds to action = 0 from the 4 matrices in the 3rd dimension of q\_values. Here we will see that (2,0) has a q value of 100, nudging the agent to move upwards (take action = 0).

```
# Individual param testing
epsilon = 0.9 # exploration
max_epsilon = 1
min_epsilon = 0.01
epsilon_decay_rate = 0.9 # If you go beyond 1 it will make epsilon go to inf
discount_factor = 0.1 #discount factor for future rewards (gamma)
learning_rate = 0.9 #the rate at which the AI agent should learn (alpha)
training_episodes = 1000
tau = 0.1 # This is temperature for boltzmann policy. Higher tau means more equal probability of taking an action. Lower tau is choosin

q_values = []
episodes = []
steps = []

#run through training episodes
for episode in range(training_episodes):
    # Define starting location. Keep it empty to randomise start location during training
    row_index, column_index = get_starting_location()
    step = 0

    #continue taking actions (i.e., moving) until we reach a terminal state
    #(i.e., until we reach the item packaging area or crash into an item storage location)
    while not is_terminal_state(row_index, column_index):

        # Choose the policy
        action_index = eps_greedy(row_index, column_index, epsilon) # Epsilon greedy policy
        # action_index = boltz_policy(row_index, column_index, tau) # Boltzmann policy
        # action_index = random_pol(row_index, column_index) # Random policy

        # # Decay the epsilon rate linearly
        # if epsilon > 0.01:
        #     epsilon *= epsilon_decay_rate

        # # Temp decay
        # if tau > 1.:
        #     tau -= 0.1

        # Decay the epsilon rate exponentially based on minimum and max values
        epsilon = min_epsilon/((max_epsilon-min_epsilon)*np.exp(-epsilon_decay_rate)) # https://medium.com/@nancyjemi/level-up-understandin
        # epsilon = min_epsilon + (max_epsilon - min_epsilon) * np.exp(-1. * episode / epsilon_decay_rate) # Non-lienar decay

        #perform the chosen action, and transition to the next state (i.e., move to the next location)
        old_row_index, old_column_index = row_index, column_index #store the old row and column indexes
        row_index, column_index = get_next_location(row_index, column_index, action_index)

        # Bellman's equation: Q[s_old,a] + alpha*(R[s_old,a] + gamma*(max(Q[s])) - Q[s_old,a])
        q_m[old_row_index, old_column_index, action_index] = q_m[old_row_index, old_column_index, action_index] + learning_rate*(rewards[ro
        print(f"\nAction taken: Go {actions[action_index]} (index = {action_index})\nMoving from {old_row_index, old_column_index} to {row_
        print(f"Q value is: {q_m[old_row_index, old_column_index, action_index]}")
        q_values.append(q_m[old_row_index, old_column_index, action_index])
        episodes.append(episode)
        step += 1
        steps.append(step)

    # print(f"\nQ matrix for going {actions[action_index]} is:\n {q_values[:, :, action_index]}")
    # #receive the reward for moving to the new state, and calculate the temporal difference
    # reward = rewards[row_index, column_index]
    # old_q_value = q_values[old_row_index, old_column_index, action_index]
    # temporal_difference = reward + (discount_factor * np.max(q_values[row_index, column_index])) - old_q_value

    # #update the Q-value for the previous state and action pair
    # new_q_value = old_q_value + (learning_rate * temporal_difference)
    # q_values[old_row_index, old_column_index, action_index] = new_q_value

print(f"\nTraining complete for {training_episodes} training episodes!")
```





Streaming output truncated to the last 5000 lines.

Moving from (0, 2) to (0, 1) with reward -1.0  
Q value is: 979.9999999999982

Action taken: Go left (index = 3)  
Moving from (0, 1) to (0, 0) with reward -1.0  
Q value is: 989.9999999999987

Action taken: Go down (index = 2)  
Moving from (0, 0) to (1, 0) with reward 100.0  
Q value is: 999.9999999999993

Action taken: Go left (index = 3)  
Moving from (7, 8) to (7, 7) with reward -1.0  
Q value is: 708.7857739213772

Action taken: Go left (index = 3)  
Moving from (7, 7) to (7, 6) with reward -1.0  
Q value is: 775.7414237570148

Action taken: Go up (index = 0)  
Moving from (7, 6) to (6, 6) with reward -1.0  
Q value is: 814.611983474748

Action taken: Go up (index = 0)  
Moving from (6, 6) to (5, 6) with reward -1.0  
Q value is: 836.2190157288333

Action taken: Go up (index = 0)  
Moving from (5, 6) to (4, 6) with reward -1.0  
Q value is: 849.2999036919308

Action taken: Go up (index = 0)  
Moving from (4, 6) to (3, 6) with reward -1.0  
Q value is: 859.8171081825539

Action taken: Go right (index = 1)  
Moving from (3, 6) to (3, 7) with reward -1.0  
Q value is: 869.9681070132146

Action taken: Go right (index = 1)  
Moving from (3, 7) to (3, 8) with reward -1.0  
Q value is: 879.995266691308

Action taken: Go up (index = 0)  
Moving from (3, 8) to (2, 8) with reward -1.0  
Q value is: 889.9994599891417

Action taken: Go up (index = 0)  
Moving from (2, 8) to (1, 8) with reward -1.0  
Q value is: 899.9998963873518

Action taken: Go up (index = 0)  
Moving from (1, 8) to (0, 8) with reward -1.0  
Q value is: 909.9999819402808

Action taken: Go left (index = 3)  
Moving from (0, 8) to (0, 7) with reward -1.0

```
results = pd.DataFrame([q_values, episodes, steps], index = ["q_values", "episodes", "steps"]).T
```

```
results_grouped_episodes = results.groupby(by = ["episodes"]).mean().reset_index()  
results_grouped_episodes
```



	episodes	q_values	steps
0	0.0	44.550000	1.5
1	1.0	-18.720000	3.0
2	2.0	-90.000000	1.0
3	3.0	-23.175000	2.5
4	4.0	-45.450000	1.5
...	...	...	...
995	995.0	277.874035	1.5
996	996.0	293.182255	4.0
997	997.0	463.842136	20.0
998	998.0	337.188428	3.5
999	999.0	-211.869125	3.5

1000 rows × 3 columns

```

import matplotlib.pyplot as plt
import seaborn as sns

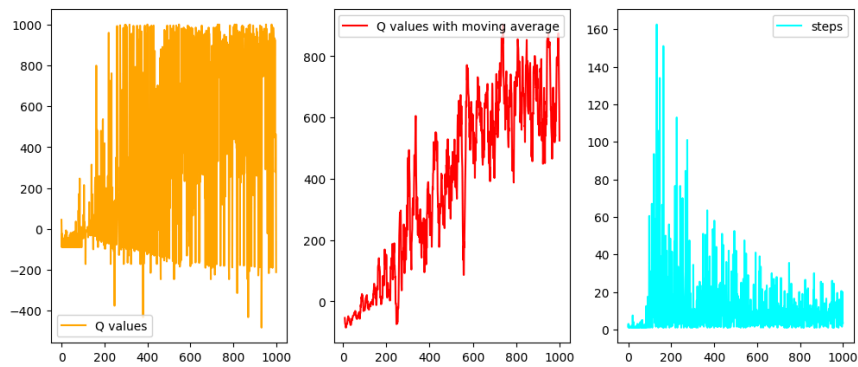
window_size = 10 # Adjust the window size as needed
rolling_avg_scores = np.convolve(results_grouped_episodes['q_values'], np.ones(window_size)/window_size, mode='valid')

fig, ax = plt.subplots(1, 3, figsize=(12, 5))
fig.suptitle("Q values for baseline agent")
ax[0].plot(results_grouped_episodes["episodes"], results_grouped_episodes['q_values'], label = "Q values", color = "orange")
ax[1].plot(range(training_episodes)[window_size - 1:], rolling_avg_scores, label = "Q values with moving average", color = "red")
ax[2].plot(results_grouped_episodes["episodes"], results_grouped_episodes['steps'], label = "steps", color = "cyan")
ax[0].legend()
ax[1].legend()
ax[2].legend()
plt.show()

```



Q values for baseline agent



## ✓ Test run

```

# epsilon
tau

```



0.1

	0	1	2	3	4	5	6	7	8	9	10
0	-1	-1	-1	-1	-1	-1	-1	-1	-1	-100	-100

```
# get_starting_location()

start = time.time()
print(f"Optimal path finding started: {datetime.fromtimestamp(start)}\n")

# Starting location
# Interesting starting locations: 9,4 | 3,8 | 5,9 | 9,10 | 1,10
# row, col = 9,10
row, col = get_starting_location(9,10)
shortest_path_result = get_shortest_path(row, col)
print()
print(f"Path: {shortest_path_result} (Moves: {len(shortest_path_result)})\n") #starting at row 3, column 9

end = time.time()
print(f"Optimal path finding ended: {datetime.fromtimestamp(end)}")
elapsed_time = end - start
print(f"Elapsed time: {np.round(elapsed_time/60, 3)} minutes\n")

# Draw path

# Define the size of the grid
n = 11
# Create a grid graph
G = nx.grid_2d_graph(n, n)
# Set the position of the nodes using the node coordinates
pos = {(x, y): (y, -x) for x, y in G.nodes()}
nx.draw(G, pos=pos, node_size=300, with_labels=False)

# Define the coordinates of the boxes you want the line to go through
# coords = [(0, 0), (1, 1), (2, 2), (3, 3), (4, 4)]
coords = [tuple(x) for x in shortest_path_result] # https://stackoverflow.com/questions/5506511/python-converting-list-of-lists-to-tuple

# Draw the line through the specified coordinates
for i in range(len(coords) - 1):
    x1, y1 = coords[i]
    x2, y2 = coords[i + 1]
    plt.plot([y1, y2], [-x1, -x2], 'r-', linewidth=2)
# Set the x and y axis ticks
# plt.xticks(range(n), [str(i) for i in range(n)]) # Column indices
# plt.yticks(range(-n, 0), [str(-i) for i in range(n)]) # Row indices
plt.show()

🔗 Optimal path finding started: 2024-05-12 14:42:25.378140
```

```

!pip install gym

!pip install plotly

# !pip install pygame
# !pip uninstall pygame

# !pip install gym[toy_text]==0.26.*
# !pip install gym[toy_text]==0.26.2
!pip install gym[box2d]==0.26.0 pygame==1.5.27 pyvirtualdisplay

!pip install pyvirtualdisplay
!pip install pygame==1.5.27

```

## ✓ Double DQN

```

# Run this

import gym
import torch
import torch.nn as nn
import torch.optim as optim
import numpy as np
import random
from collections import deque
import plotly.graph_objects as go
import matplotlib.pyplot as plt

random_seed = 42
random.seed(random_seed)
np.random.seed(random_seed)
torch.manual_seed(random_seed)

# https://www.bing.com/search?pglt=43&q=nn.Module&cvid=3e8744507f7d432fb59a2a95abeccad8&gs\_lcrp=EgZjaHJvbnUyBggAEUEUYOTIGCAEQABhAMgYIAhAAGEAyBggDEAAYQDIGCAQQABhAMgYIBRAAGEAyBggGEAAYQDIGCAcQABhAMgYICBBFGDzSAQcxNDhqMGoxqAIAsAIA&FORM=A

# Define Double DQN Model Architecture
class DoubleDQN(nn.Module):
    def __init__(self, state_size, action_size, hidden_sizes):
        super(DoubleDQN, self).__init__()
        self.hidden1 = nn.Linear(state_size, hidden_sizes[0])
        self.hidden2 = nn.Linear(hidden_sizes[0], hidden_sizes[1])
        self.output = nn.Linear(hidden_sizes[1], action_size)

    def forward(self, state):
        x = torch.relu(self.hidden1(state))
        x = torch.relu(self.hidden2(x))
        return self.output(x)

# Function to visualize Q-values predicted by the policy network
def visualize_q_values(policy_net, state):
    with torch.no_grad():
        q_values = policy_net(torch.tensor(state, dtype=torch.float32).unsqueeze(0))
        action_values = q_values.numpy()[0]
        actions = np.arange(len(action_values))

    plt.bar(actions, action_values)
    plt.title("Optimal Q-values for each action")
    plt.ylabel("Q-value")
    plt.xlabel("Action")
    plt.show()

    # fig = go.Figure(data=[go.Bar(x=actions, y=action_values)])
    # fig.update_layout(title='Q-values Predicted by Policy Network',
    #                    xaxis_title='Action',
    #                    yaxis_title='Q-value')
    # fig.show()

# # Function to compare parameters of policy network and target network
# def compare_networks(policy_net, target_net):
#     policy_params = np.concatenate([param.data.numpy().flatten() for param in policy_net.parameters()])
#     target_params = np.concatenate([param.data.numpy().flatten() for param in target_net.parameters()])

#     fig = go.Figure()
#     fig.add_trace(go.Scatter(y=policy_params, mode='lines', name='Policy Network', line=dict(color='blue', width=2), opacity=1))
#     fig.add_trace(go.Scatter(y=target_params, mode='lines', name='Target Network', line=dict(color='red', width=2), opacity=0.5))

```

```

#     fig.add_trace(go.Scatter(y=target_params, mode= lines , name= target_network , line=dict(color= red , width=2), opacity=0.5))

#     # Add vertical lines to indicate the separation between different layers
#     layer_sizes = [state_size] + hidden_sizes + [action_size]
#     param_count = 0
#     for size in layer_sizes[:-1]:
#         param_count += size * layer_sizes[layer_sizes.index(size) + 1]
#         fig.add_shape(type='line',
#             x0=param_count, y0=policy_params.min(),
#             x1=param_count, y1=policy_params.max(),
#             line=dict(color='gray', width=1))

#     fig.update_layout(title='Comparison of Policy and Target Networks Parameters',
#         xaxis_title='Parameter Index',
#         yaxis_title='Parameter Value',
#         legend=dict(x=0.8, y=0.9, bgcolor='rgba(255, 255, 255, 0.8)'))
#     fig.show()

# Set up the environment
env = gym.make('MountainCar-v0')
state_size = env.observation_space.shape[0]
action_size = env.action_space.n

# Reward Function
def reward_function(state, next_state):
    position = state[0]
    next_position = next_state[0]
    velocity = state[1]
    next_velocity = next_state[1]

    # Check if the episode is done
    if next_position >= 0.5:
        return 100 # Large positive reward for reaching the target position

    # Reward proportional to the change in position towards the target
    reward = (next_position - position) * 10

    # Additional reward for maintaining positive velocity
    if next_velocity > 0:
        reward += 1

    return reward

# Hyperparameters
batch_size = 64
gamma = 0.8
epsilon_start = 1.0
epsilon_end = 0.01
epsilon_decay = 0.95
# epsilon = 0.9 # Only run if using linear epsilon decay
target_update_freq = 10
learning_rate = 0.01
memory_size = 10000
hidden_sizes = [100, 64]
num_episodes = 100
print_every = 10

# Create the Double DQN network and target network
policy_net = DoubleDQN(state_size, action_size, hidden_sizes)
target_net = DoubleDQN(state_size, action_size, hidden_sizes)
target_net.load_state_dict(policy_net.state_dict())
target_net.eval()

# Define the optimizer
optimizer = optim.Adam(policy_net.parameters(), lr=learning_rate)

# Define the replay memory
# https://docs.python.org/3/library/collections.html
memory = deque(maxlen=memory_size)

# Define the epsilon-greedy policy
def epsilon_greedy_policy(state, epsilon):
    if random.random() > epsilon:
        with torch.no_grad():
            return policy_net(state).argmax(dim=1).item()
    else:
        return random.randrange(action_size)

# Initialize lists to store scores and episode numbers
scores = []

```

```

episode_numbers = []

# Train the Double DQN agent
for episode in range(num_episodes):
    state = env.reset()

    # Epsilon decay strategy
    epsilon = epsilon_end + (epsilon_start - epsilon_end) * np.exp(-1. * episode / epsilon_decay) # Non-linear decay
    # epsilon *= epsilon_decay # Linear decay

    done = False
    score = 0
    while not done:
        action = epsilon_greedy_policy(torch.tensor(state, dtype=torch.float32).unsqueeze(0), epsilon)
        next_state, _, done, _ = env.step(action)
        reward = reward_function(state, next_state)
        memory.append((state, action, reward, next_state, done))
        state = next_state
        score += reward

    if len(memory) >= batch_size:
        batch = random.sample(memory, batch_size)
        states, actions, rewards, next_states, dones = zip(*batch)
        states = torch.tensor(states, dtype=torch.float32)
        actions = torch.tensor(actions, dtype=torch.long).unsqueeze(1)
        rewards = torch.tensor(rewards, dtype=torch.float32).unsqueeze(1)
        next_states = torch.tensor(next_states, dtype=torch.float32)
        dones = torch.tensor(dones, dtype=torch.float32).unsqueeze(1)

        q_values = policy_net(states).gather(1, actions)
        next_q_values = target_net(next_states).max(1)[0].unsqueeze(1)
        expected_q_values = rewards + (1 - dones) * gamma * next_q_values

        loss = nn.MSELoss()(q_values, expected_q_values)
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    if episode % target_update_freq == 0:
        target_net.load_state_dict(policy_net.state_dict())

    scores.append(score)
    episode_numbers.append(episode)

    # Print episode number, epsilon, and score after every 'print_every' episodes
    if (episode + 1) % print_every == 0:
        print(f"Episode: {episode+1}, Epsilon: {epsilon:.4f}, Score: {score:.2f}")

# Visualize Q-values for a sample state
sample_state = env.observation_space.sample()
visualize_q_values(policy_net, sample_state)

# # Compare parameters of policy network and target network
# compare_networks(policy_net, target_net)

# # Plot the scores against episode numbers using Plotly
fig = go.Figure(data=go.Scatter(x=episode_numbers, y=scores, mode='lines'))
fig.update_layout(title='Double DQN - Score vs Episode',
#                 xaxis_title='Episode',
#                 yaxis_title='Score')
fig.show()

# Calculate rolling average
window_size = 10 # Adjust the window size as needed
rolling_avg_scores = np.convolve(scores, np.ones(window_size)/window_size, mode='valid')

# # Plot the scores and rolling average against episode numbers using Plotly
fig = go.Figure()
fig.add_trace(go.Scatter(x=episode_numbers, y=scores, mode='lines', name='Scores'))
fig.add_trace(go.Scatter(x=episode_numbers[window_size - 1:], y=rolling_avg_scores, mode='lines', name='Rolling Average'))
fig.update_layout(title='Double DQN - Score vs Episode with Rolling Average',
#                 xaxis_title='Episode',
#                 yaxis_title='Score',
#                 legend=dict(x=0.8, y=0.9, bgcolor='rgba(255, 255, 255, 0.8)'))
fig.show()

fig, ax = plt.subplots(1, 2)
ax[0].plot(episode_numbers, scores, label = "Scores", color = "orange")
ax[1].plot(episode_numbers[window_size - 1:], rolling_avg_scores, label = "Rolling Average", color = "cyan")
ax[0].legend()
ax[1].legend()
plt.title("Double DQN - Score vs Episode with Rolling Average")
plt.show()

```

```
.
# import plotly.graph_objects as go

# fig = go.Figure(layout=dict(title="Double DQN - Score vs Episode with Rolling Average"))

# # Plot scores
# trace1 = go.Scatter(
#     x=episode_numbers,
#     y=scores,
#     mode="lines",
#     name="Scores",
#     line_color="orange",
# )

# # Plot rolling average scores (starting from window_size)
# trace2 = go.Scatter(
#     x=episode_numbers[window_size - 1:],
#     y=rolling_avg_scores,
#     mode="lines",
#     name="Rolling Average",
#     line_color="cyan",
# )



# # Add traces to the figure
# fig.add_trace(trace1)
# fig.add_trace(trace2)

# # Show legend
# fig.update_layout(legend=dict(yanchor="top", y=1.02, xanchor="right", x=1))

# fig.show()

# Close the environment
env.close()
```

Episode: 40, Epsilon: 0.0100, Score: 99.18  
Episode: 50, Epsilon: 0.0100, Score: 116.47  
Episode: 60, Epsilon: 0.0100, Score: 121.37  
Episode: 70, Epsilon: 0.0100, Score: 116.98  
Episode: 80, Epsilon: 0.0100, Score: 118.99  
Episode: 90, Epsilon: 0.0100, Score: 116.83  
Episode: 100, Epsilon: 0.0100, Score: 130.93



✓ PER DQN



```

import gym
import torch
import torch.nn as nn
import torch.optim as optim
import numpy as np
import random
from collections import namedtuple
import matplotlib.pyplot as plt
import plotly.graph_objects as go

random_seed = 42
random.seed(random_seed)
np.random.seed(random_seed)
torch.manual_seed(random_seed)

# https://pytorch.org/r1/stable/reference/generated/torchrl.data.PrioritizedReplayBuffer.html
# https://www.bing.com/search?pglt=43&q=nn.Module&cvid=3e8744507f7d432fb59a2a95abeccad8&gs_lcrp=
# EgZjaHJvbWUyBggAEUEYOTIGCAEQABhAMgYIAhAAGEAyBggDEAAYQDIGCAQABhAMgYIBRAAGEAyBggGEAAAYQDIGCacQABhAMgYICBBFGDzSAQcxNDhqMGoxqAIAAsAIA&FORI

# Define the Prioritized Experience Replay buffer
class PrioritizedReplayBuffer:
    def __init__(self, capacity, alpha=0.6, beta_start=0.4, beta_frames=100000):
        self.capacity = capacity
        self.alpha = alpha
        self.beta_start = beta_start
        self.beta_frames = beta_frames
        self.buffer = []
        self.pos = 0
        self.priorities = np.zeros((capacity,), dtype=np.float32)

    def add(self, state, action, reward, next_state, done):
        max_prio = self.priorities.max() if self.buffer else 1.0
        if len(self.buffer) < self.capacity:
            self.buffer.append(None)
        self.buffer[self.pos] = (state, action, reward, next_state, done)
        self.priorities[self.pos] = max_prio
        self.pos = (self.pos + 1) % self.capacity

    def sample(self, batch_size, beta):
        if len(self.buffer) == self.capacity:
            prios = self.priorities
        else:
            prios = self.priorities[:self.pos]
        probs = prios ** self.alpha
        probs /= probs.sum()
        indices = np.random.choice(len(self.buffer), batch_size, p=probs)
        samples = [self.buffer[idx] for idx in indices]
        total = len(self.buffer)
        weights = (total * probs[indices]) ** (-beta)
        weights /= weights.max()
        return samples, indices, np.array(weights, dtype=np.float32)

    def update_priorities(self, batch_indices, batch_priorities):
        for idx, prio in zip(batch_indices, batch_priorities):
            self.priorities[idx] = prio

    def __len__(self):
        return len(self.buffer)

# Define the DQN network architecture
class DQN(nn.Module):
    def __init__(self, state_size, action_size, hidden_sizes):
        super(DQN, self).__init__()
        self.hidden1 = nn.Linear(state_size, hidden_sizes[0])
        self.hidden2 = nn.Linear(hidden_sizes[0], hidden_sizes[1])
        # self.hidden3 = nn.Linear(hidden_sizes[1], hidden_sizes[2])
        # self.hidden4 = nn.Linear(hidden_sizes[2], hidden_sizes[3])
        # self.hidden5 = nn.Linear(hidden_sizes[3], hidden_sizes[4])
        self.output = nn.Linear(hidden_sizes[1], action_size)

    def forward(self, state):
        x = torch.relu(self.hidden1(state))
        x = torch.relu(self.hidden2(x))
        # x = torch.relu(self.hidden3(x))
        # x = torch.relu(self.hidden4(x))
        # x = torch.relu(self.hidden5(x))
        return self.output(x)

# Reward Function
def reward_function(state, next_state):
    position = state[0]
    next_position = next_state[0]

```



```

velocity = state[1]
next_velocity = next_state[1]

# Check if the episode is done
if next_position >= 0.5:
    return 100 # Large positive reward for reaching the target position

# Reward proportional to the change in position towards the target
reward = (next_position - position) * 10

# Additional reward for maintaining positive velocity
if next_velocity > 0:
    reward += 1

return reward

# Function to visualize Q-values predicted by the policy network
def visualize_q_values(policy_net, state):
    with torch.no_grad():
        q_values = policy_net(torch.tensor(state, dtype=torch.float32).unsqueeze(0))
        action_values = q_values.numpy()[0]
        actions = np.arange(len(action_values))
        plt.bar(actions, action_values)
        plt.title("Optimal Q-values for each action")
        plt.ylabel("Q-value")
        plt.xlabel("Action")

    # fig = go.Figure(data=[go.Bar(x=actions, y=action_values)])
    # fig.update_layout(title='Q-values Predicted by Policy Network',
    #                    xaxis_title='Action',
    #                    yaxis_title='Q-value')
    # fig.show()

# # Function to compare parameters of policy network and target network
# def compare_networks(policy_net, target_net):
#     policy_params = np.concatenate([param.data.numpy().flatten() for param in policy_net.parameters()])
#     target_params = np.concatenate([param.data.numpy().flatten() for param in target_net.parameters()])

#     fig = go.Figure()
#     fig.add_trace(go.Scatter(y=policy_params, mode='lines', name='Policy Network', line=dict(color='rgba(0, 0, 255, 0.8)', width=2)))
#     fig.add_trace(go.Scatter(y=target_params, mode='lines', name='Target Network', line=dict(color='rgba(255, 0, 0, 0.6)', width=2)))
#     fig = go.Figure()
#     fig.add_trace(go.Scatter(y=policy_params, mode='lines', name='Policy Network', line=dict(color='blue', width=2, opacity=1)))
#     fig.add_trace(go.Scatter(y=target_params, mode='lines', name='Target Network', line=dict(color='red', width=2, opacity=0.3)))

#     # Add vertical lines to indicate the separation between different layers
#     layer_sizes = [state_size] + hidden_sizes + [action_size]
#     param_count = 0
#     for size in layer_sizes[:-1]:
#         param_count += size * layer_sizes[layer_sizes.index(size) + 1]
#         fig.add_shape(type='line',
#                      x0=param_count, y0=policy_params.min(),
#                      x1=param_count, y1=policy_params.max(),
#                      line=dict(color='gray', width=1, dash='dot'))

#     fig.update_layout(title='Comparison of Policy and Target Networks Parameters',
#                      xaxis_title='Parameter Index',
#                      yaxis_title='Parameter Value',
#                      legend=dict(x=0.8, y=0.9, bgcolor='rgba(255, 255, 255, 0.8)'))
#     fig.show()

# Set up the environment
env = gym.make('MountainCar-v0')
state_size = env.observation_space.shape[0]
action_size = env.action_space.n

# Hyperparameters
batch_size = 64
gamma = 0.8
learning_rate = 0.0001
target_update_freq = 10
epsilon_start = 1.0
epsilon_end = 0.01
epsilon_decay = 0.85
# epsilon = 0.9 # Use for linear decay rate
memory_size = 10000
alpha = 0.4
beta_start = 0.6
beta_frames = 50000
hidden_sizes = [100, 64]
num_episodes = 100

```

```

print_every = 10

# Create the PER buffer and DQN network
memory = PrioritizedReplayBuffer(memory_size, alpha)
policy_net = DQN(state_size, action_size, hidden_sizes)
target_net = DQN(state_size, action_size, hidden_sizes)
target_net.load_state_dict(policy_net.state_dict())
target_net.eval()

# Define the optimizer
optimizer = optim.Adam(policy_net.parameters(), lr=learning_rate)

# Define the loss function
def loss_fn(batch_states, batch_actions, batch_rewards, batch_next_states, batch_dones, batch_weights):
    batch_q_values = policy_net(batch_states).gather(1, batch_actions)
    batch_next_q_values = target_net(batch_next_states).max(1)[0].unsqueeze(1).detach()
    batch_expected_q_values = batch_rewards + (1 - batch_dones) * gamma * batch_next_q_values
    batch_weights_tensor = torch.FloatTensor(batch_weights).unsqueeze(1)
    loss = (batch_weights_tensor * (batch_q_values - batch_expected_q_values) ** 2).mean()
    return loss, batch_q_values, batch_expected_q_values

# Training loop
scores = []
episode_numbers = []
for episode in range(num_episodes):
    state = env.reset()
    score = 0
    done = False
    beta = min(1.0, beta_start + episode * (1.0 - beta_start) / beta_frames)

    # Epsilon decay rate strategy
    epsilon = epsilon_end + (epsilon_start - epsilon_end) * np.exp(-1. * episode / epsilon_decay) # Non-linear decay
    # epsilon *= epsilon_decay # Linear decay

    while not done:
        if random.random() < epsilon:
            action = env.action_space.sample()
        else:
            with torch.no_grad():
                state_tensor = torch.FloatTensor(state).unsqueeze(0)
                q_values = policy_net(state_tensor)
                action = q_values.argmax().item()

        next_state, _, done, _ = env.step(action)
        reward = reward_function(state, next_state)
        memory.add(state, action, reward, next_state, done)
        state = next_state
        score += reward

        if len(memory) >= batch_size:
            batch, batch_indices, batch_weights = memory.sample(batch_size, beta)
            batch_states, batch_actions, batch_rewards, batch_next_states, batch_dones = zip(*batch)
            batch_states = torch.FloatTensor(batch_states)
            batch_actions = torch.LongTensor(batch_actions).unsqueeze(1)
            batch_rewards = torch.FloatTensor(batch_rewards).unsqueeze(1)
            batch_next_states = torch.FloatTensor(batch_next_states)
            batch_dones = torch.FloatTensor(batch_dones).unsqueeze(1)

            loss, batch_q_values, batch_expected_q_values = loss_fn(batch_states, batch_actions, batch_rewards, batch_next_states, batch_dones, batch_weights)
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

            batch_priorities = (torch.abs(batch_q_values - batch_expected_q_values) + 1e-5).squeeze().detach().cpu().numpy()
            memory.update_priorities(batch_indices, batch_priorities)

    if episode % target_update_freq == 0:
        target_net.load_state_dict(policy_net.state_dict())

    scores.append(score)
    episode_numbers.append(episode)

    if (episode + 1) % print_every == 0:
        avg_score = np.mean(scores[-print_every:])
        print(f"Episode: {episode+1}, Average Score: {avg_score:.2f}, Epsilon: {epsilon:.2f}")

# Visualize Q-values for a sample state
sample_state = env.observation_space.sample()

```

✓ Referenced from: [https://github.com/radical-p/Breakout\\_Deep\\_Q\\_Reinforcement\\_Learning/blob/main/Breakout\\_DQN.py](https://github.com/radical-p/Breakout_Deep_Q_Reinforcement_Learning/blob/main/Breakout_DQN.py)

```
# !pip install stable-baselines3[extra] pyvirtualdisplay gym[atari] pygame
# !pip install atari-py==0.2.5
# !apt-get install unrar

# !wget http://www.atarimania.com/roms/Roms.rar
# !unrar x Roms.rar
# !unzip ROMS.zip
# !python -m atari_py.import_roms ROMS

# import os

# os.system('apt-get update')
# os.system('apt-get install -y xvfb')
# os.system('wget https://raw.githubusercontent.com/yandexdataschool/Practical_DL/fall18/xvfb -O ../xvfb')
# os.system('apt-get install -y python-opengl ffmpeg')
# os.system('pip install pygame==1.5.0')

# os.system('python -m pip install -U pygame --user')

# prefix = 'https://raw.githubusercontent.com/yandexdataschool/Practical_RL/master/week04_approx_rl/'

# os.system('wget ' + prefix + 'atari_wrappers.py')
# os.system('wget ' + prefix + 'utils.py')
# os.system('wget ' + prefix + 'replay_buffer.py')
# os.system('wget ' + prefix + 'framebuffer.py')

# # print('setup complete')

# # Xvfb will be launched if you run on a server
# import os
# if type(os.environ.get("DISPLAY")) is not str or len(os.environ.get("DISPLAY")) == 0:
#     !bash ../xvfb start
#     os.environ['DISPLAY'] = ':1'

import random
import numpy as np
import pandas as pd
import gym
import time
from collections import deque
from keras import optimizers
from keras.models import Sequential
from keras.layers import Dense
import matplotlib.pyplot as plt
import tensorflow as tf

class DQN:
    def __init__(self, env):
        # Initialize the DQN agent
        self.env = env
        self.memory = deque(maxlen=4000) # Replay memory (Original: 400000)
        self.gamma = 0.9 # Discount factor for future rewards
        self.epsilon = .9 # Exploration rate
        self.epsilon_min = 0.01 # Minimum exploration rate
        self.epsilon_decay = self.epsilon_min / 200000 # Decay rate for exploration

        self.batch_size = 16 # Batch size for training
        self.train_start = 100 # Number of experiences required before starting training
        self.state_size = self.env.observation_space.shape[0] * 4 # Size of the state vector
        self.action_size = self.env.action_space.n # Number of possible actions
        self.learning_rate = 0.0001 # Learning rate for the optimizer

        self.evaluation_model = self.create_model() # Neural network for evaluation
        self.target_model = self.create_model() # Neural network as a target for stable training

    def create_model(self):
        # Create a neural network model for the DQN
        model = Sequential()
        model.add(Dense(128 * 2, input_dim=self.state_size, activation='relu'))
        model.add(Dense(128 * 2, activation='relu'))
        # model.add(Dense(128 * 2, activation='relu'))
        model.add(Dense(self.action_size, activation='linear'))
        model.compile(loss='mean_squared_error', optimizer=tf.keras.optimizers.keras.optimizers.RMSprop(lr=self.learning_rate, decay=0.99, epsilon=
```

```

    return model

def choose_action(self, state, steps):
    # Choose an action using epsilon-greedy exploration strategy
    if steps > 50000:
        if self.epsilon > self.epsilon_min:
            self.epsilon -= self.epsilon_decay
    if np.random.random() < self.epsilon:
        return self.env.action_space.sample()
    return np.argmax(self.evaluation_model.predict(state)[0])

def remember(self, cur_state, action, reward, new_state, done):
    # Store the experience in the replay memory
    if not hasattr(self, 'memory_counter'):
        self.memory_counter = 0
    transition = (cur_state, action, reward, new_state, done)
    self.memory.extend([transition])
    self.memory_counter += 1

def replay(self):
    # Train the DQN by replaying experiences from the replay memory
    if len(self.memory) < self.train_start:
        return
    mini_batch = random.sample(self.memory, self.batch_size)
    update_input = np.zeros((self.batch_size, self.state_size))
    update_target = np.zeros((self.batch_size, self.action_size))

    for i in range(self.batch_size):
        state, action, reward, new_state, done = mini_batch[i]
        target = self.evaluation_model.predict(state)[0]
        if done:
            target[action] = reward
        else:
            target[action] = reward + self.gamma * np.amax(self.target_model.predict(new_state)[0])
        update_input[i] = state
        update_target[i] = target

    self.evaluation_model.fit(update_input, update_target, batch_size=self.batch_size, epochs=1, verbose=0)

def target_train(self):
    # Update the target model with the weights of the evaluation model
    self.target_model.set_weights(self.evaluation_model.get_weights())
    return

def visualize(self, reward, episode):
    # Visualize the average reward per episode
    plt.plot(episode, reward, 'ob-')
    plt.title('Average reward vs episode')
    plt.ylabel('Reward')
    plt.xlabel('Episodes')
    plt.grid()
    plt.show()

def transform(self, state):
    # Transform the state representation if necessary
    if state.shape[1] == 512:
        return state
    a = [np.binary_repr(x, width=8) for x in state[0]]
    res = []
    for x in a:
        res.extend([x[:2], x[2:4], x[4:6], x[6:]])
    res = [int(x, 2) for x in res]
    return np.array(res)

def main():
    # Initialize the environment
    env = gym.make('Breakout-ram-v0')
    env = env.unwrapped

    # Print environment information
    print(env.action_space)
    print(env.observation_space.shape[0])
    print(env.observation_space.high)
    print(env.observation_space.low)

    episodes = 250
    trial_len = 10

    tmp_reward = 0
    sum_rewards = 0
    total_steps = 0

```

```

graph_reward = []
graph_episodes = []
time_record = []

dqn_agent = DQN(env=env)

for i_episode in range(episodes):
    start_time = time.time()
    total_reward = 0
    cur_state_tuple = env.reset()
    cur_state = np.array(cur_state_tuple[0]).reshape(1, 128)
    cur_state = dqn_agent.transform(cur_state).reshape(1, 128 * 4) / 4

    for _ in range(trial_len):
        # Choose action, take a step in the environment
        action = dqn_agent.choose_action(cur_state, total_steps)
        step_result = env.step(action)
        if len(step_result) == 5:
            new_state, reward, done, _, info = step_result
        else:
            new_state, reward, done, _, info = step_result + (None,) * (4 - len(step_result))

        new_state = new_state.reshape(1, 128)
        new_state = dqn_agent.transform(new_state).reshape(1, 128 * 4) / 4
        total_reward += reward
        sum_rewards += reward
        tmp_reward += reward
        if reward > 0:
            reward = 1 # Testing whether it is good.

        # Store the experience in the replay memory
        dqn_agent.remember(cur_state, action, reward, new_state, done)

        if total_steps > 100:
            if total_steps % 4 == 0:
                # Train the DQN by replaying experiences from the replay memory
                dqn_agent.replay()
            if total_steps % 500 == 0:
                # Update the target model with the weights of the evaluation model
                dqn_agent.target_train()

        cur_state = new_state
        total_steps += 1
        if done:
            env.reset()
            break

    if (i_episode + 1) % 5 == 0:
        graph_reward.append(sum_rewards / 5)
        graph_episodes.append(i_episode + 1)
        sum_rewards = 0

    end_time = time.time()
    time_record.append(end_time - start_time)
    tmp_reward = 0

print("Reward: ")
print(graph_reward)
print("Episode: ")
print(graph_episodes)
print("Average_time: ")
print(sum(time_record) / 60)
dqn_agent.visualize(graph_reward, graph_episodes)

```

```
main()
```

```

/usr/local/lib/python3.10/dist-packages/gym/envs/registration.py:555: UserWarning
  logger.warn(
Discrete(4)
128
[255 255 255 255 255 255 255 255 255 255 255 255 255 255 255 255 255
 255 255 255 255 255 255 255 255 255 255 255 255 255 255 255 255
 255 255 255 255 255 255 255 255 255 255 255 255 255 255 255 255
 255 255 255 255 255 255 255 255 255 255 255 255 255 255 255 255
 255 255 255 255 255 255 255 255 255 255 255 255 255 255 255 255
 255 255 255 255 255 255 255 255 255 255 255 255 255 255 255 255
 255 255]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
/usr/local/lib/python3.10/dist-packages/keras/src/optimizers/legacy/rmsprop.py:14:
  super().__init__(name, **kwargs)
Streaming output truncated to the last 5000 lines.
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 43ms/step
1/1 [=====] - 0s 48ms/step
1/1 [=====] - 0s 41ms/step
1/1 [=====] - 0s 42ms/step
1/1 [=====] - 0s 52ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 51ms/step
1/1 [=====] - 0s 37ms/step
1/1 [=====] - 0s 48ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 48ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 56ms/step
1/1 [=====] - 0s 44ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 53ms/step
1/1 [=====] - 0s 69ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 41ms/step
1/1 [=====] - 0s 52ms/step
1/1 [=====] - 0s 51ms/step
```



## ✓ Set up

Code taken from: [https://pytorch.org/tutorials/intermediate/reinforcement\\_ppo.html?highlight=ppo](https://pytorch.org/tutorials/intermediate/reinforcement_ppo.html?highlight=ppo)

Another article used for reference: <https://medium.com/aureliantactics/ppo-hyperparameters-and-ranges-6fc2d29bccbe>

```
# !pip3 install torchrl
# !pip3 install gym[mujoco]
# !pip3 install tqdm

from collections import defaultdict

import matplotlib.pyplot as plt
import torch
from tensordict.nn import TensorDictModule
from tensordict.nn.distributions import NormalParamExtractor
from torch import nn
from torchrl.collectors import SyncDataCollector
from torchrl.data.replay_buffers import ReplayBuffer
from torchrl.data.replay_buffers.samplers import SamplerWithoutReplacement
from torchrl.data.replay_buffers.storages import LazyTensorStorage
from torchrl.envs import (Compose, DoubleToFloat, ObservationNorm, StepCounter,
                          TransformedEnv)
from torchrl.envs.libs.gym import GymEnv
from torchrl.envs.utils import check_env_specs, ExplorationType, set_exploration_type
from torchrl.modules import ProbabilisticActor, TanhNormal, ValueOperator
from torchrl.objectives import ClipPPOLoss
from torchrl.objectives.value import GAE
from tqdm import tqdm

import multiprocessing
is_fork = multiprocessing.get_start_method() == "fork"
device = (
    torch.device(0)
    if torch.cuda.is_available() and not is_fork
    else torch.device("cpu")
)
num_cells = 256 # number of cells in each layer i.e. output dim.
lr = 0.0005
max_grad_norm = 1.0

/usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283: DeprecationWarning: `should_run_async` will not call `transform_
and should_run_async(code)

frames_per_batch = 1000
# For a complete training, bring the number of frames up to 1M
total_frames = 100_000

sub_batch_size = 64 # cardinality of the sub-samples gathered from the current data in the inner loop
num_epochs = 10 # optimization steps per batch of data collected

# Clipping Range: 0.1, 0.2, 0.3
clip_epsilon = (
    0.2 # clip value for PPO loss
)

# Can be thought of as bias-variance trade off
gamma = 0.9 # Discount factor
lmbda = 0.1
entropy_eps = 0.9 # Regularizer (0 to 0.01)

base_env = GymEnv("InvertedDoublePendulum-v4", device=device)

env = TransformedEnv(
    base_env,
    Compose(
        # normalize observations
        ObservationNorm(in_keys=["observation"]),
        DoubleToFloat(),
        StepCounter(),
    ),
)
```


```
 /usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283: DeprecationWarning: `should_run_async` will not call `transform_` and `should_run_async` (code)
```

```
env.transform[0].init_stats(num_iter=1000, reduce_dim=0, cat_dim=0)
```

```
print("normalization constant shape:", env.transform[0].loc.shape)
print("observation_spec:", env.observation_spec)
print("reward_spec:", env.reward_spec)
print("input_spec:", env.input_spec)
print("action_spec (as defined by input_spec):", env.action_spec)
```

```
 normalization constant shape: torch.Size([11])
observation_spec: CompositeSpec(
  observation: UnboundedContinuousTensorSpec(
    shape=torch.Size([11]),
    space=None,
    device=cpu,
    dtype=torch.float32,
    domain=continuous),
  step_count: BoundedTensorSpec(
    shape=torch.Size([1]),
    space=ContinuousBox(
      low=Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.int64, contiguous=True),
      high=Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.int64, contiguous=True)),
    device=cpu,
    dtype=torch.int64,
    domain=continuous), device=cpu, shape=torch.Size([]))
reward_spec: UnboundedContinuousTensorSpec(
  shape=torch.Size([1]),
  space=ContinuousBox(
    low=Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.float32, contiguous=True),
    high=Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.float32, contiguous=True)),
  device=cpu,
  dtype=torch.float32,
  domain=continuous)
input_spec: CompositeSpec(
  full_state_spec: CompositeSpec(
    step_count: BoundedTensorSpec(
      shape=torch.Size([1]),
      space=ContinuousBox(
        low=Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.int64, contiguous=True),
        high=Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.int64, contiguous=True)),
      device=cpu,
      dtype=torch.int64,
      domain=continuous), device=cpu, shape=torch.Size([])),
    full_action_spec: CompositeSpec(
      action: BoundedTensorSpec(
        shape=torch.Size([1]),
        space=ContinuousBox(
          low=Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.float32, contiguous=True),
          high=Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.float32, contiguous=True)),
        device=cpu,
        dtype=torch.float32,
        domain=continuous), device=cpu, shape=torch.Size([]), device=cpu, shape=torch.Size([]))
  action_spec (as defined by input_spec): BoundedTensorSpec(
    shape=torch.Size([1]),
    space=ContinuousBox(
      low=Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.float32, contiguous=True),
      high=Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.float32, contiguous=True)),
    device=cpu,
    dtype=torch.float32,
    domain=continuous)
```

```
rollout = env.rollout(3)
print("rollout of three steps:", rollout)
print("Shape of the rollout TensorDict:", rollout.batch_size)
```

```
 rollout of three steps: TensorDict(
  fields={
    action: Tensor(shape=torch.Size([3, 1]), device=cpu, dtype=torch.float32, is_shared=False),
    done: Tensor(shape=torch.Size([3, 1]), device=cpu, dtype=torch.bool, is_shared=False),
    next: TensorDict(
      fields={
        done: Tensor(shape=torch.Size([3, 1]), device=cpu, dtype=torch.bool, is_shared=False),
        observation: Tensor(shape=torch.Size([3, 11]), device=cpu, dtype=torch.float32, is_shared=False),
        reward: Tensor(shape=torch.Size([3, 1]), device=cpu, dtype=torch.float32, is_shared=False),
        step_count: Tensor(shape=torch.Size([3, 1]), device=cpu, dtype=torch.int64, is_shared=False),
        terminated: Tensor(shape=torch.Size([3, 1]), device=cpu, dtype=torch.bool, is_shared=False),
        truncated: Tensor(shape=torch.Size([3, 1]), device=cpu, dtype=torch.bool, is_shared=False),
        batch_size: torch.Size([3]),
        device: cpu,
        is_shared: False),
    observation: Tensor(shape=torch.Size([3, 11]), device=cpu, dtype=torch.float32, is_shared=False),
    step_count: Tensor(shape=torch.Size([3, 1]), device=cpu, dtype=torch.int64, is_shared=False),
    terminated: Tensor(shape=torch.Size([3, 1]), device=cpu, dtype=torch.bool, is_shared=False),
```

```

        truncated: Tensor(shape=torch.Size([3, 1]), device=cpu, dtype=torch.bool, is_shared=False)),
        batch_size=torch.Size([3]),
        device=cpu,
        is_shared=False)
Shape of the rollout TensorDict: torch.Size([3])

```

## ✚ Policy

```

actor_net = nn.Sequential(
    nn.LazyLinear(num_cells, device=device),
    nn.Tanh(),
    nn.LazyLinear(num_cells, device=device),
    nn.Tanh(),
    nn.LazyLinear(num_cells, device=device),
    nn.Tanh(),
    nn.LazyLinear(2 * env.action_spec.shape[-1], device=device),
    NormalParamExtractor(),
)

```

 /usr/local/lib/python3.10/dist-packages/torch/nn/modules/lazy.py:181: UserWarning: Lazy modules are a new feature under heavy development
warnings.warn('Lazy modules are a new feature under heavy development '

```

policy_module = TensorDictModule(
    actor_net, in_keys=["observation"], out_keys=["loc", "scale"]
)

```

```

policy_module = ProbabilisticActor(
    module=policy_module,
    spec=env.action_spec,
    in_keys=["loc", "scale"],
    distribution_class=TanhNormal,
    distribution_kwargs={
        "min": env.action_spec.space.low,
        "max": env.action_spec.space.high,
    },
    return_log_prob=True,
    # we'll need the log-prob for the numerator of the importance weights
)

```

```

value_net = nn.Sequential(
    nn.LazyLinear(num_cells, device=device),
    nn.Tanh(),
    nn.LazyLinear(num_cells, device=device),
    nn.Tanh(),
    nn.LazyLinear(num_cells, device=device),
    nn.Tanh(),
    nn.LazyLinear(1, device=device),
)

```

```


value_module = ValueOperator(
    module=value_net,
    in_keys=["observation"],
)

```

```

print("Running policy:", policy_module(env.reset()))
print("Running value:", value_module(env.reset()))

```

 Running policy: TensorDict(
 fields={
 action: Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.float32, is\_shared=False),
 done: Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.bool, is\_shared=False),
 loc: Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.float32, is\_shared=False),
 observation: Tensor(shape=torch.Size([11]), device=cpu, dtype=torch.float32, is\_shared=False),
 sample\_log\_prob: Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.float32, is\_shared=False),
 scale: Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.float32, is\_shared=False),
 step\_count: Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.int64, is\_shared=False),
 terminated: Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.bool, is\_shared=False),
 truncated: Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.bool, is\_shared=False)},
 batch\_size=torch.Size([1]),
 device=cpu,
 is\_shared=False)
Running value: TensorDict(
 fields={
 done: Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.bool, is\_shared=False),
 observation: Tensor(shape=torch.Size([11]), device=cpu, dtype=torch.float32, is\_shared=False),
 state\_value: Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.float32, is\_shared=False),
 step\_count: Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.int64, is\_shared=False),
 terminated: Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.bool, is\_shared=False),

```

        truncated: Tensor(shape=torch.Size([1]), device=cpu, dtype=torch.bool, is_shared=False)),
        batch_size=torch.Size([]),
        device=cpu,
        is_shared=False)

collector = SyncDataCollector(
    env,
    policy_module,
    frames_per_batch=frames_per_batch,
    total_frames=total_frames,
    split_trajs=False,
    device=device,
)

replay_buffer = ReplayBuffer(
    storage=LazyTensorStorage(max_size=frames_per_batch),
    sampler=SamplerWithoutReplacement(),
)

# Generalized Advantage Estimation
advantage_module = GAE(
    gamma=gamma, lmbda=lmbda, value_network=value_module, average_gae=True
)

loss_module = ClipPPOLoss(
    actor_network=policy_module,
    critic_network=value_module,
    clip_epsilon=clip_epsilon,
    entropy_bonus=bool(entropy_eps),
    entropy_coef=entropy_eps,
    # these keys match by default but we set this for completeness
    critic_coef=1.0,
    loss_critic_type="smooth_l1",
)

optim = torch.optim.Adam(loss_module.parameters(), lr)
scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(
    optim, total_frames // frames_per_batch, 0.0
)

```

```

logs = defaultdict(list)
pbar = tqdm(total=total_frames)
eval_str = ""


# We iterate over the collector until it reaches the total number of frames it was
# designed to collect:
for i, tensordict_data in enumerate(collector):
    # we now have a batch of data to work with. Let's learn something from it.
    for _ in range(num_epochs):
        # We'll need an "advantage" signal to make PPO work.
        # We re-compute it at each epoch as its value depends on the value
        # network which is updated in the inner loop.
        advantage_module(tensordict_data)
        data_view = tensordict_data.reshape(-1)
        replay_buffer.extend(data_view.cpu())
    for _ in range(frames_per_batch // sub_batch_size):
        subdata = replay_buffer.sample(sub_batch_size)
        loss_vals = loss_module(subdata.to(device))
        loss_value = (
            loss_vals["loss_objective"]
            + loss_vals["loss_critic"]
            + loss_vals["loss_entropy"]
        )

        # Optimization: backward, grad clipping and optimization step
        loss_value.backward()
        # this is not strictly mandatory but it's good practice to keep
        # your gradient norm bounded
        torch.nn.utils.clip_grad_norm_(loss_module.parameters(), max_grad_norm)
        optim.step()
        optim.zero_grad()

    logs["reward"].append(tensordict_data["next", "reward"].mean().item())
    pbar.update(tensordict_data.numel())
    cum_reward_str = (
        f"average reward={logs['reward'][-1]: 4.4f} (init={logs['reward'][0]: 4.4f})"
    )
    logs["step_count"].append(tensordict_data["step_count"].max().item())
    stepcount_str = f"step count (max): {logs['step_count'][-1]}"
    logs["lr"].append(optim.param_groups[0]["lr"])
    lr_str = f"lr policy: {logs['lr'][-1]: 4.4f}"
    if i % 10 == 0:
        # We evaluate the policy once every 10 batches of data.
        # Evaluation is rather simple: execute the policy without exploration
        # (take the expected value of the action distribution) for a given
        # number of steps (1000, which is our ``env`` horizon).
        # The ``rollout`` method of the ``env`` can take a policy as argument:
        # it will then execute this policy at each step.
        with set_exploration_type(ExplorationType.MEAN), torch.no_grad():
            # execute a rollout with the trained policy
            eval_rollout = env.rollout(1000, policy_module)
            logs["eval reward"].append(eval_rollout["next", "reward"].mean().item())
            logs["eval reward (sum)"].append(
                eval_rollout["next", "reward"].sum().item()
            )
            logs["eval step_count"].append(eval_rollout["step_count"].max().item())
            eval_str = (
                f"eval cumulative reward: {logs['eval reward (sum)'][-1]: 4.4f} "
                f"(init: {logs['eval reward (sum)'][0]: 4.4f}), "
                f"eval step-count: {logs['eval step_count'][-1]}"
            )
            del eval_rollout
        pbar.set_description("", ".join([eval_str, cum_reward_str, stepcount_str, lr_str]))

# We're also using a learning rate scheduler. Like the gradient clipping,
# this is a nice-to-have but nothing necessary for PPO to work.
scheduler.step()

```

 /usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283: DeprecationWarning: `should\_run\_async` will not call `transform\_` and `should\_run\_async(code)`

eval cumulative reward: 278.8475 (init: 101.3024), eval step-count: 29, average reward= 9.2570 (init= 9.0916), step count (max):



```
plt.figure(figsize=(10, 10))
plt.subplot(2, 2, 1)
plt.plot(logs["reward"])
plt.title("training rewards (average)")
plt.subplot(2, 2, 2)
plt.plot(logs["step_count"])
plt.title("Max step count (training)")
plt.subplot(2, 2, 3)
plt.plot(logs["eval reward (sum)"])
plt.title("Return (test)")
plt.subplot(2, 2, 4)
plt.plot(logs["eval step_count"])
plt.title("Max step count (test)")
plt.show()
```

